

Table 1: CIFAR-10 posterior sampling results for CNF prior. We report expected classifier log probability, FID scores, and ELBO for the class posteriors, averaged over all 10 classes.

Model	Sampler	$\mathbb{E}[\log p(\mathbf{y} \mid \mathbf{x})]$ ( $\uparrow$ )	FID ( $\downarrow$ )	ELBO ( $\uparrow$ )
I-CFM	Prior	-5.88	84.79	-24.04
	DPS	-2.22	84.96	-
	RTB	-4.20	90.77	-147.69
	Latent HMC	-2.80	46.69	-
	Adj. Matching	-3.09	19.45	-17.23
	<b>Outsourced Diff.</b>	-3.35	34.28	-20.36

Table 2: SD 1.5 fine-tuning results, averaged across three prompts used in [Venkatraman et al.]<sup>2</sup>. DPOK, DDPO and RTB results taken from the same paper.

Sampler	$\mathbb{E}[\log r(\mathbf{x}, \mathbf{y})]$ ( $\uparrow$ )	CLIP diversity ( $\uparrow$ )
Prior	-0.17	0.18
DDPO	1.37	0.09
DPOK	1.23	0.13
RTB	1.4	0.11
<b>Outsourced Diff.</b>	1.26	0.14

Table 3: SD3 prior and posterior results with ELBO for each prompt

Prompt	Prior			Outsourced Diff.		
	Reward	Diversity	ELBO	Reward	Diversity	ELBO
A cat and a dog.	0.5	0.14	3.55	1.23	0.09	27.25
A cat riding a llama.	0.79	0.18	1.01	1.53	0.14	10.83
A quiet village is disrupted by a meteor strike.	0.65	0.24	1.29	0.94	0.21	23.2
A human with a horse face and a human with a wolf face.	1.22	0.2	19.32	1.36	0.18	26.1
<b>AVG</b>	<b>0.79</b>	<b>0.19</b>	<b>6.29</b>	<b>1.27</b>	<b>0.16</b>	<b>21.85</b>

<sup>2</sup>'A green rabbit.', 'A cat and a dog.', and 'Four roses.'

---

**Algorithm 1** Training loop for Outsourced Diffusion Sampler

---

- 1: **Initialize:** deterministic prior function  $f$  (e.g. CNF integrators, GAN generator), randomly initialized noise posterior model  $p_F^\phi$ , randomly initialized  $Z^\phi(\mathbf{y})$  (scalar for fixed  $\mathbf{y}$ ), VP-SDE backward policy  $p_B$ , log reward function  $\log r(\mathbf{x}, \mathbf{y})$ , on-policy update fraction  $p$ .
  - 2: **for** each step  $n = 1, 2, \dots, N$  **do**
  - 3:   Sample a batch of trajectories:  $\{\tau^{(i)}\}_{i=1}^B \sim p_F^\phi(\tau \mid \mathbf{y})$
  - 4:   **for**  $i = 1, \dots, B$  **do**
  - 5:     Compute log density:  $\log R^{(i)} \leftarrow \log \mathcal{N}(\mathbf{z}^{(i)}; \mathbf{0}, \mathbf{I}) + r(f(\mathbf{z}^{(i)}), \mathbf{y})$
  - 6:     Store experience  $(\tau^{(i)}, \log R^{(i)})$  in replay buffer  $\mathcal{D}$
  - 7:   **end for**
  - 8:   Draw  $u \sim \text{Uniform}(0, 1)$
  - 9:   **if**  $u \leq p$  **then**
  - 10:     Keep on-policy batch  $\{(\tau^{(i)}, \log R^{(i)})\}_{i=1}^B$
  - 11:   **else**
  - 12:     Sample off-policy batch  $\{(\tau^{(i)}, \log R^{(i)})\}_{i=1}^B \sim \mathcal{D}$
  - 13:   **end if**
  - 14:   Compute  $\mathcal{L}_{\text{TB}}(\tau; \mathbf{y}, \phi)$  for batch using TB loss eq(4).
  - 15:   Update  $p_F^\phi$ ,  $Z^\phi(\mathbf{y})$  using  $\nabla_\phi \mathcal{L}_{\text{TB}}(\tau; \mathbf{y}, \phi)$ .
  - 16: **end for**
-